

CS251 Fall 2025
(cs251.stanford.edu)



Ethereum: mechanics

Dan Boneh

Note: HW#2 posted. Due Oct. 21.

Consensus (SMR): quick summary

Goals: Safety ($LOG_t^i \leq LOG_s^j$ or $LOG_s^j \leq LOG_t^i$) and Liveness (no censorship)

Network models: synchronous vs. asynchronous

Settings:

- Open participation: need sybil resistance (PoW or PoS)
- We saw two types of consensus protocols:
 - **Nakamoto-style consensus:** longest (heaviest) chain fork choice rule
 - dynamic availability, but no finality (no safety when asynchronous)
 - **PBFT-style** (block finalized when $\geq t$ votes): finality and accountable safety
- Ethereum proof-of-stake: a finality chain that is a prefix of an available chain

New topic: limitations of Bitcoin

Recall: UTXO contains (hash of) ScriptPK

- simple script: indicates conditions when UTXO can be spent

Limitations:

- Difficult to maintain state in multi-stage contracts
- Difficult to enforce global rules on assets

A simple example: rate limiting. My wallet manages 100 UTXOs.

- Desired policy: can only transfer 2BTC per day out of my wallet

An example: DNS

Domain name system on the blockchain: [google.com → IP addr]

Need support for three operations:

- **Name.new(OwnerAddr, DomainName)**: intent to register
- **Name.update(DomainName, newVal, newOwner, OwnerSig)**
- **Name.lookup(DomainName)**

Note: also need to ensure no front-running on **Name.new()**

A broken implementation

Name.new() and Name.update() create a UTXO with ScriptPK:

**DUP HASH256 <OwnerAddr> EQVERIFY CHECKSIG VERIFY
<DNS> <DomainName> <IPAddr> <1>**

only owner can “spend” this UTXO to update domain data

Contract: (should be enforced by miners)

if domain google.com is registered,
no one else can register that domain

verify
sig is valid

ensure top
of stack is 1

Problem: this contract cannot be enforced using Bitcoin script

What to do?

NameCoin: a fork of Bitcoin that implements this contract
(see also the Ethereum Name Service -- ENS)

Can we build a blockchain that natively supports generic contracts like this?

⇒ Ethereum

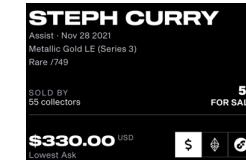


Ethereum: enables a world of applications

A world of Ethereum Decentralized apps (DAPPs)

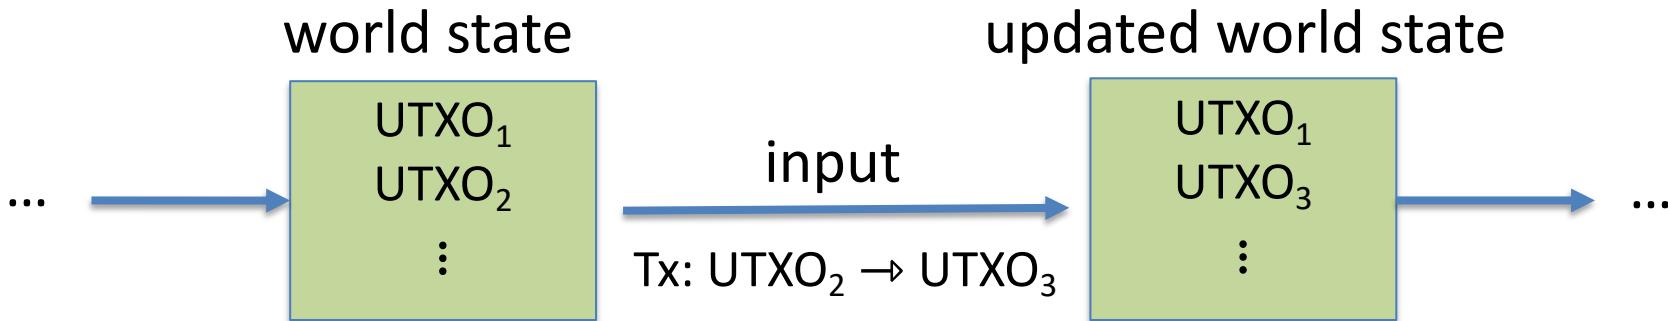
- New coins: ERC-20 standard interface
- **DeFi**: exchanges, lending, stablecoins, derivatives, etc.
- **Insurance**
- **DAOs**: decentralized organizations
- **NFTs**: Managing asset ownership (ERC-721 interface)

■
■
■
■



dappradar.com/rankings/protocol/ethereum

Bitcoin as a state transition system



Bitcoin rules:

$$F_{\text{bitcoin}} : S \times I \rightarrow S$$

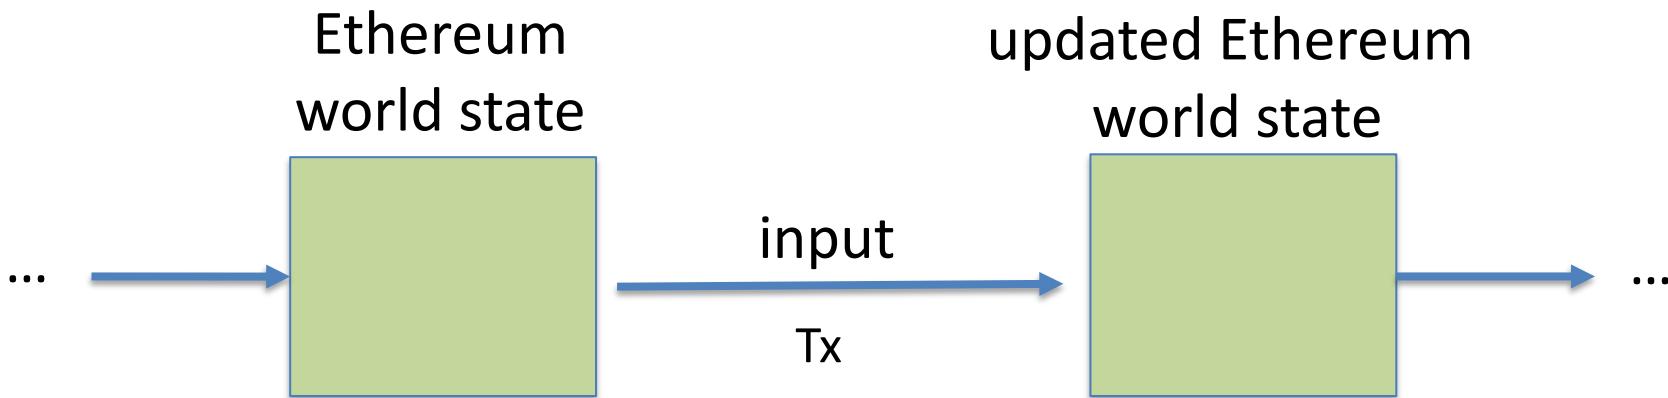
S: set of all possible world states, $s_0 \in S$ genesis state

I: set of all possible inputs

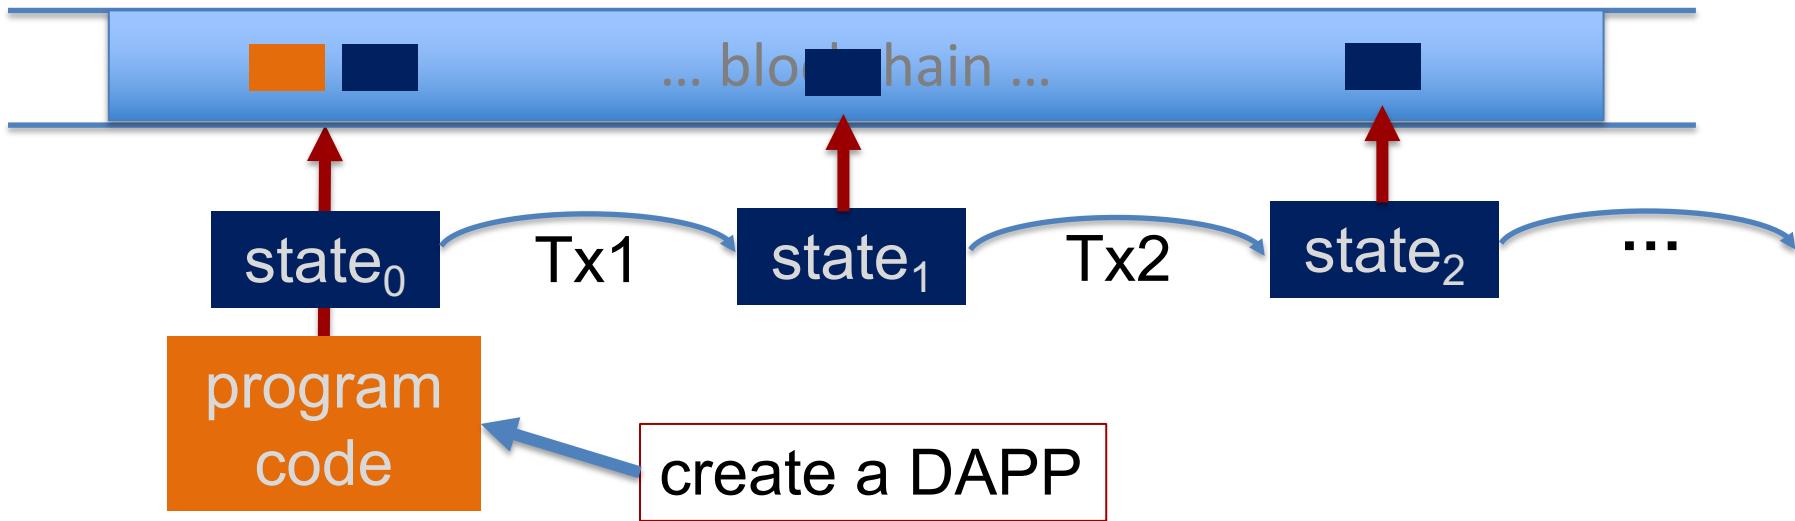
Ethereum as a state transition system

Much richer state transition functions

⇒ one transition executes an entire program



Running a program on a blockchain (DAPP)



compute layer (execution chain): The EVM

consensus layer (beacon chain)

The Ethereum system

Proof-of-Stake consensus

Block	Slot	Age	Blobs	Txn	Fee Recipient
23570698	12796465	12 secs ago	6 (67%)	283	Titan Builder
23570697	12796464	24 secs ago	3 (33%)	219	Titan Builder
23570696	12796463	36 secs ago	9 (100%)	285	Fee Recipient: 0xe68...127
23570695	12796462	48 secs ago	4 (44%)	68	beaverbuild
23570694	12796461	1 min ago	3 (33%)	117	Lido: Execution Layer Rew...
23570693	12796460	1 min ago	6 (67%)	206	Titan Builder
23570692	12796459	1 min ago	3 (33%)	278	Titan Builder
23570691	12796458	1 min ago	0 (0%)	187	BuilderNet

source: etherscan.io

One slot every 12 seconds.
(about 200 Tx per block)

One block proposer chosen
for each slot (from validator set)

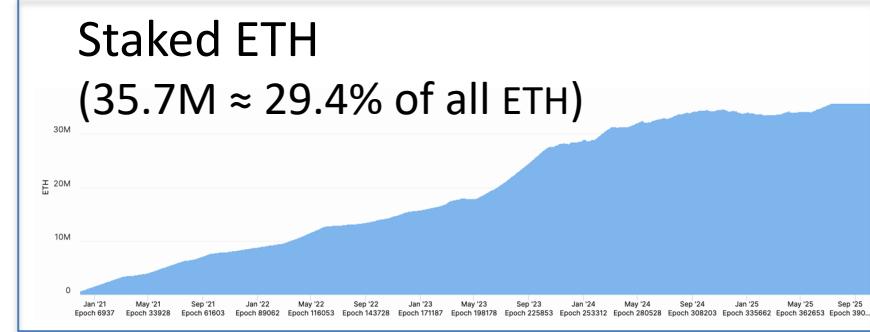
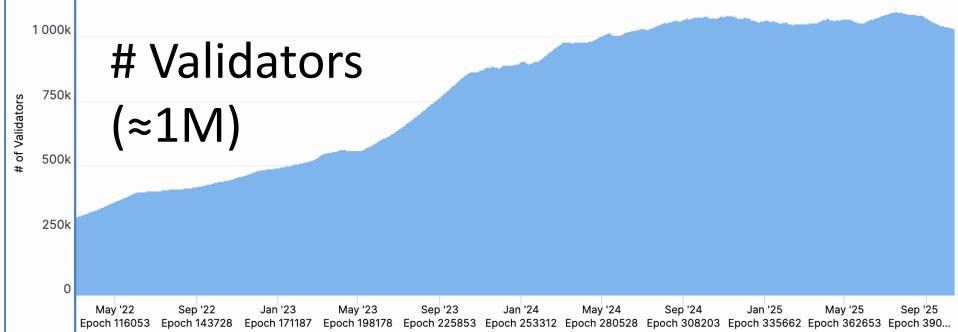
If it sends a valid block,
receives Tx fees for block
(along with other rewards)

A bit about the beacon chain (Eth2 consensus layer)

To become a validator: stake (lock up) 32 ETH ... or use Lido.

Validators:

- sign blocks to express correctness (finalized once enough sigs)
- occasionally act as ***block proposer*** (chosen at random)
- correct behavior \Rightarrow issued **new ETH** every epoch (32 blocks)
- incorrect behavior \Rightarrow slashed (lots of details)



The economics of staking

Validator locks up 32 ETH. Oct 2025: 35.7M ETH staked (total)

Annual validator income (an example):

- Sources: issuance, Tx fees (tips), MEV
- Total: 0.93 ETH/year (2.93% on 32 ETH staked)

In practice: staking provider (e.g., Lido) takes a cut of the returns

An Ethereum node

update
world state

compute layer (execution chain)

`notify_new_payload(payload)` [Engine API]

sends transactions to compute layer



32 slots
in an epoch

(12sec × 32 = 6.4 mins)

consensus layer (beacon chain)

The Ethereum Compute Layer: The EVM

Ethereum compute layer: the EVM

World state: set of accounts identified by 20-byte address.

Two types of accounts:

(1) externally owned accounts (EOA):

controlled by ECDSA signing key pair (pk,sk).

sk: signing key known only to account owner

Since May 2025:
EOA can also be
controlled by code

(2) contracts: controlled by code.

code set at account creation time, does not change

Data associated with an account

Account data

Owned (EOA)

Contracts

address (computed): $H(pk)$ $H(\text{CreatorAddr}, \text{CreatorNonce})$

code: \perp (or address) CodeHash

storage root (state): \perp StorageRoot

balance (in Wei): balance balance (1 Wei = 10^{-18} ETH)

nonce: nonce nonce

(#Tx sent) + (#accounts created): anti-replay mechanism

Account state: persistent storage

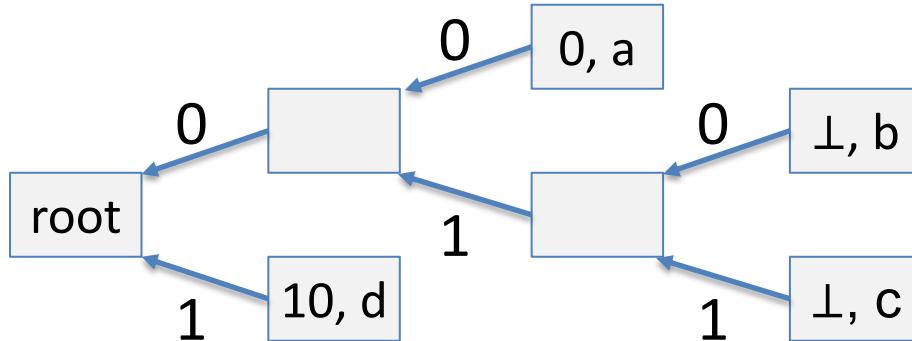
Every contract has an associated **storage array $S[]$** :

$S[0], S[1], \dots, S[2^{256}-1]$: each cell holds 32 bytes, init to 0.

Account storage root: **Merkle Patricia Tree hash of $S[]$**

- Cannot compute full Merkle tree hash: 2^{256} leaves

$S[000] = a$
 $S[010] = b$
 $S[011] = c$
 $S[110] = d$



time to compute root hash:
 $\leq 2 \times |S|$

$|S| = \# \text{ non-zero cells}$

State transitions: Tx and messages

Transactions: signed data by initiator

- **To:** 32-byte address of target ($0 \rightarrow$ create new account)
- **From, [Signature]:** initiator address and signature on Tx (if owned)
- **Value:** # Wei being sent with Tx (1 Wei = 10^{-18} ETH)
- Tx fees (EIP 1559): **gasLimit, maxFee, maxPriorityFee** (later)
- if $To = 0$: create new contract **code = (init, body)**
- if $To \neq 0$: **data** (what function to call & arguments)
- **nonce**: must match current nonce of sender (prevents Tx replay)
- **chain_id**: ensures Tx can only be submitted to the intended chain

State transitions: Tx and messages

Transaction types:

owned → owned: transfer ETH between users

owned → contract: call contract with ETH & data

Example (block #10993504)

<u>From</u>	<u>To</u>	<u>msg.value</u>	<u>Tx fee (ETH)</u>
0xa4ec1125ce9428ae5...	0x2cebe81fe0dcd220e...	0 Ether	0.00404405
0xba272f30459a119b2...	Uniswap V2: Router 2	0.14 Ether	0.00644563
0x4299d864bbda0fe32...	Uniswap V2: Router 2	89.839104111882671 Ether	0.00716578
0x4d1317a2a98cfea41...	0xc59f33af5f4a7c8647...	14.501 Ether	0.001239
0x29ecaa773f052d14e...	CryptoKitties: Core	0 Ether	0.00775543
0x63bb46461696416fa...	Uniswap V2: Router 2	0.203036474328481 Ether	0.00766728
0xde70238aef7a35abd...	Balancer: ETH/DOUGH...	0 Ether	0.00261582
0x69aca10fe1394d535f...	0x837d03aa7fc09b8be...	0 Ether	0.00259936
0xe2f5d180626d29e75...	Uniswap V2: Router 2	0 Ether	0.00665809

Messages: virtual Tx initiated by a contract

Same as Tx, but no signature (contract has no signing key)

contract → owned: contract sends funds to user

contract → contract: one program calls another (and sends funds)

One Tx from user: can lead to many Tx processed. Composability!

Tx from owned addr → contract → another contract



another contract → different owned

Example Tx

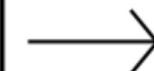
State

14c5f8ba: owned
- 1024 eth

bb75a980: contract
- 5202 eth
if !contract.storage[tx.data[0]]:
contract.storage[tx.data[0]] = tx.data[1]
[0, 235235, 0, ALICE]

892bf92f: contract
- 0 eth
send(tx.value / 3, contract.storage[0])
send(tx.value / 3, contract.storage[1])
send(tx.value / 3, contract.storage[2])
[ALICE, BOB, CHARLIE]

4096ad65: owned
- 77 eth



Transaction

From: 14c5f8ba
To: bb75a980
Value: 10 eth
Data: 2,
CHARLIE
Sig: 30452fdedb3d
f7959f2ceb8a1



State'

14c5f8ba:
- 1014 eth

bb75a980:
- 5212 eth
if !contract.storage[tx.data[0]]:
contract.storage[tx.data[0]] = tx.data[1]
[0, 235235, CHARLIE, ALICE ..]

892bf92f:
- 0 eth
send(tx.value / 3, contract.storage[0])
send(tx.value / 3, contract.storage[1])
send(tx.value / 3, contract.storage[2])
[ALICE, BOB, CHARLIE]

4096ad65:
- 77 eth

world state (four accounts)

updated world state

An Ethereum Block

Block proposer creates a block of n Tx: (from Txs submitted by users)

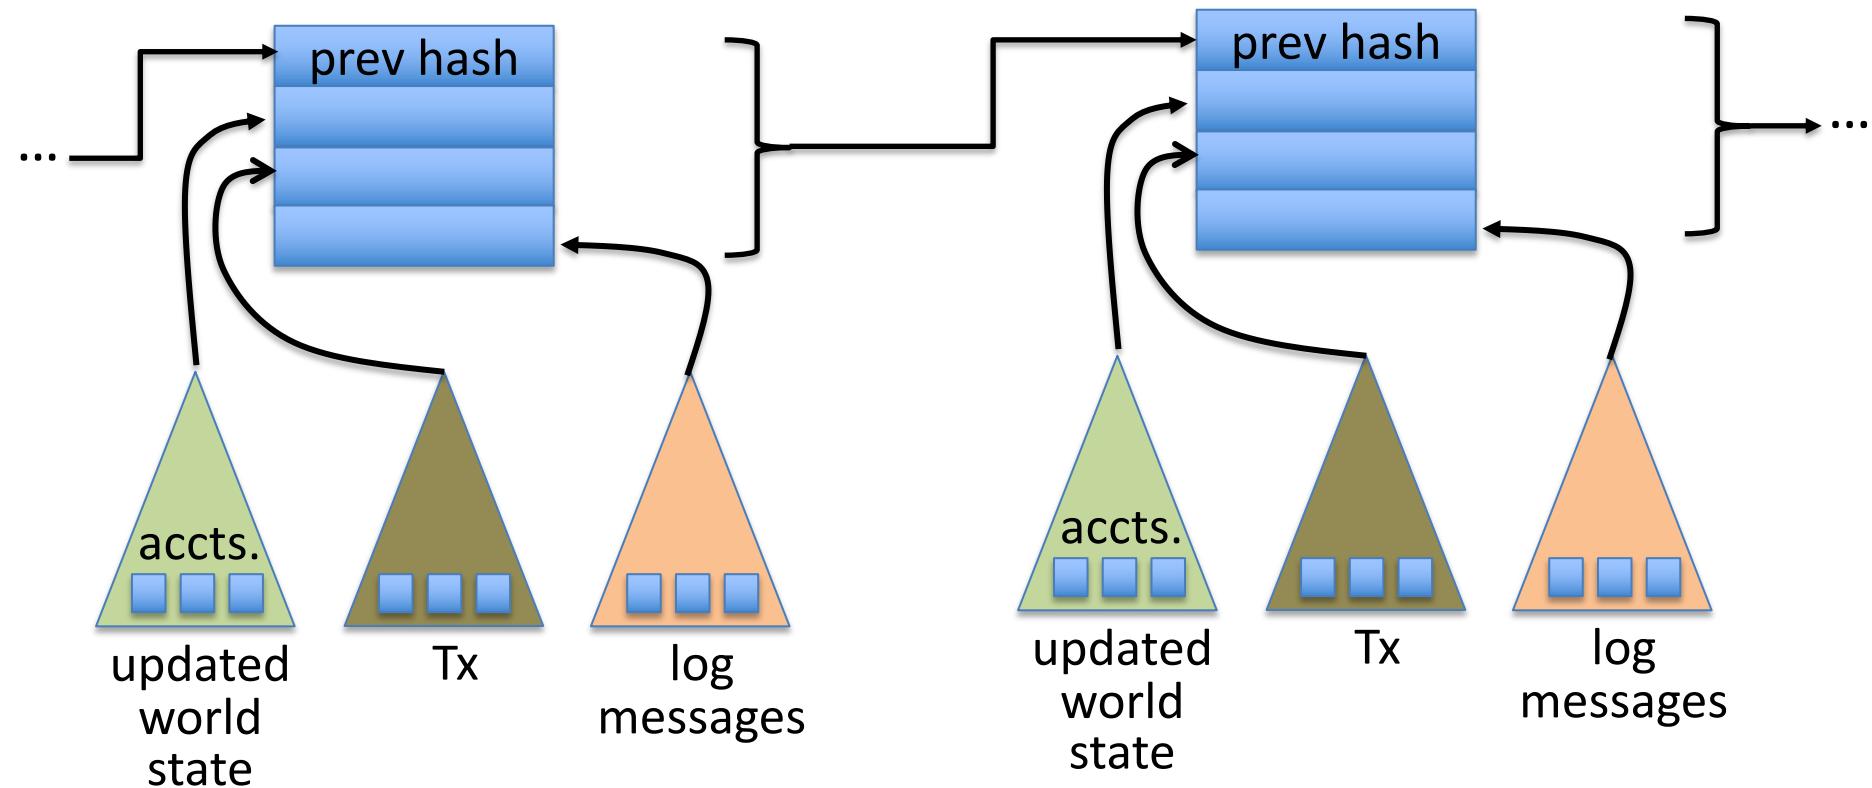
- To produce a block do:
 - for $i=1, \dots, n$: execute state change of Tx_i sequentially
(can change state of $>n$ accounts)
 - record updated world state in block

Other validators re-execute all Tx to verify block \Rightarrow
sign block if valid \Rightarrow enough sigs, epoch is finalized.

Block header data (simplified)

- (1) consensus data: proposer ID, parent hash, votes, etc.
- (2) address of gas beneficiary: where Tx fees will go
- (3) world state root:** updated world state
 - Merkle Patricia Tree hash of all accounts in the system
- (4) **Tx root:** Merkle hash of all Tx processed in block
- (5) **Tx receipt root:** Merkle hash of log messages generated in block
- (6) **Gas Used:** used to adjust the gas price (max gas per block is 45M)

The Ethereum blockchain: abstractly



Amount of memory to run a node



ETH total blockchain size (archival): 24 TB (Oct. 2025)

An example contract: NameCoin

```
contract nameCoin {      // Solidity code (next lecture)

    struct nameEntry {
        address owner;    // address of domain owner
        bytes32 value;    // IP address
    }

    // array of all registered domains
    mapping (bytes32 => nameEntry) data;
```

An example contract: NameCoin

```
function nameNew(bytes32 name) {  
    // registration costs is 100 Wei  
  
    if (data[name] == 0 && msg.value >= 100) {  
        data[name].owner = msg.sender // record domain owner  
        emit Register(msg.sender, name) // log event  
    }  
}
```



Code ensures that no one can take over a registered name

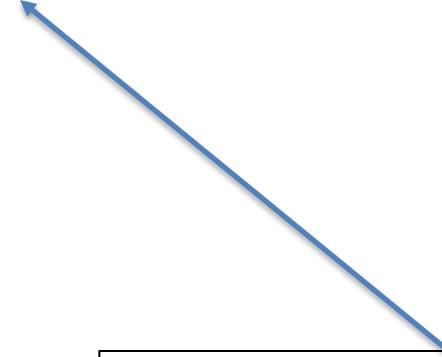
Serious bug in this code! Front running. Solved using commitments.

An example contract: NameCoin

```
function nameUpdate(  
    bytes32 name, bytes32 newValue, address newOwner) {  
  
    // check if message is from domain owner,  
    // and update cost of 10 Wei is paid  
  
    if (data[name].owner == msg.sender && msg.value >= 10) {  
        data[name].value = newValue;          // record new value  
        data[name].owner = newOwner;         // record new owner  
    }}}
```

An example contract: NameCoin

```
function nameLookup(bytes32 name) {  
    return data[name];  
}  
  
} // end of contract
```



Used by other contracts
Humans do not need this
(use etherscan.io)

EVM mechanics: execution environment

Write code in Solidity (or another front-end language)

- ⇒ compile to EVM bytecode, e.g., using **solc** compiler
(some projects use WASM or BPF bytecode)
- ⇒ validators use the EVM to execute contract bytecode
in response to a Tx

The EVM

Stack machine (like Bitcoin) but with JUMP

- max stack depth = 1024
- program aborts if stack size exceeded; block proposer keeps gas
- a contract can create or call another contract
 - There are several ways to call another contract
 - Using CALL to call a contract that is different from the caller creates a new volatile execution frame that is deleted on return.

The EVM memory types

The EVM has three types of zero initialized memory per contract.

All three are private to the contract that owns them (e.g., nameCoin)

- **Persistent storage** (on blockchain): SLOAD, SSTORE (expensive)
- **Volatile memory** (for a single Tx, one per execution frame): MLOAD, MSTORE (very cheap, 3 Gas)
- **Transient memory** (for a single Tx, but behaves like storage): TLOAD, TSTORE (cheap, 100 Gas)
- LOG0(data): write data to log

Every instruction costs gas, examples:

SSTORE **addr** (32 bytes), **value** (32 bytes)

- zero → non-zero: 20,000 gas
- non-zero → non-zero: 5,000 gas (for a cold slot)
- non-zero → zero: 15,000 gas refund (example)

Refund is given for reducing size of blockchain state

CREATE : 32,000 + 200 × (code size) gas;

CALL **maxgas**, **addr**, **value**, args

SELFDESTRUCT **addr**: kill current contract (5000 gas + 25K gas if **addr** is empty)

Gas calculation

Why charge gas?

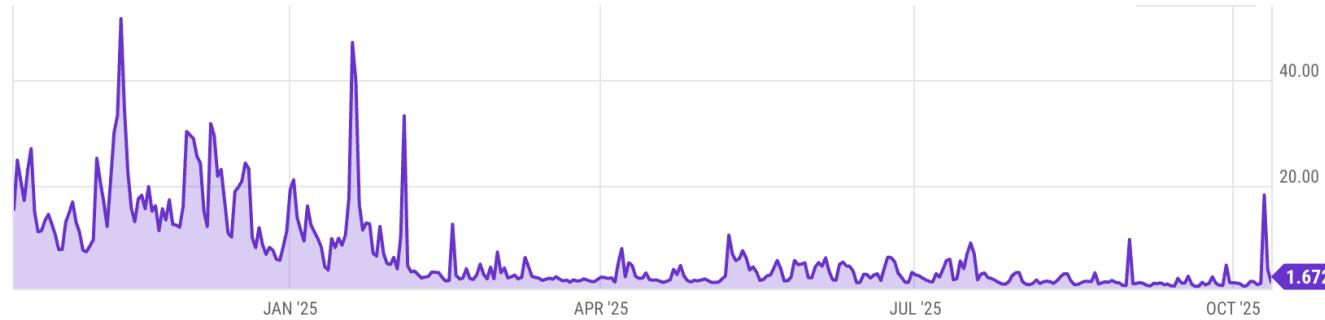
- Tx fees (gas) prevents submitting Tx that runs for many steps.
- During high load: block proposer chooses Tx from mempool that maximize its income.

Old EVM: (prior to EIP1559, live on 8/2021)

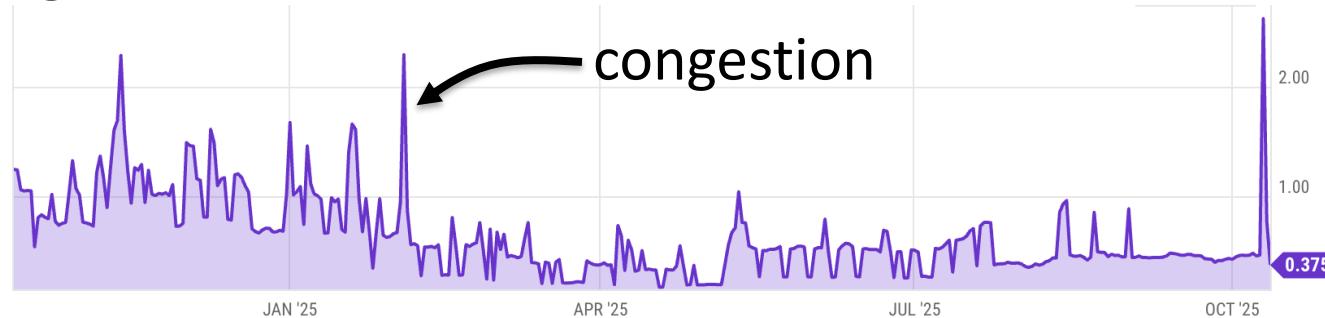
- Every Tx contains a gasPrice ``bid'' (gas \rightarrow Wei conversion price)
- Producer chooses Tx with highest gasPrice (max $\text{sum}(\text{gasPrice} \times \text{gasLimit})$)
 \Rightarrow not an efficient auction mechanism (first price auction)

Gas prices spike during congestion

GasPrice in Gwei: $1.672 \text{ Gwei} = 1.672 \times 10^{-9} \text{ ETH}$



Average Tx fee in USD:



Gas calculation: EIP1559 (since 8/2021)

EIP1559 goals (informal):

- users incentivized to bid their true utility for posting Tx,
- block proposer incentivized to not create fake Tx, and
- disincentivize off chain agreements.

[Transaction Fee Mechanism Design, by T. Roughgarden, 2021]

Gas calculation: EIP1559

Every block has a “baseFee”:

the **minimum** gasPrice for all Tx in the block

baseFee is computed from total gas in earlier blocks:

- earlier blocks at gas limit (45M gas) \Rightarrow base fee goes up 12.5%
- earlier blocks empty \Rightarrow base fee decreases by 12.5%

interpolate
in between

If earlier blocks at “target size” (22.5M gas) \Rightarrow base fee does not change

Gas calculation

EIP1559 Tx specifies three parameters:

- **gasLimit**: max total gas allowed for Tx
- **maxFee**: maximum allowed gas price (max gas → Wei conversion)
- **maxPriorityFee**: additional “tip” to be paid to block proposer

Computed **gasPrice** bid:

```
gasPrice ← min(maxFee, baseFee + maxPriorityFee)
```

Max Tx fee: **gasLimit** × **gasPrice**

Gas calculation (informal)

gasUsed \leftarrow gas used by Tx

Send $\text{gasUsed} \times (\text{gasPrice} - \text{baseFee})$ to block proposer

BURN $\text{gasUsed} \times \text{baseFee}$



\Rightarrow total supply of ETH can decrease

Gas calculation

- (1) if **gasPrice** < **baseFee**: abort
 - (2) If **gasLimit** × **gasPrice** < msg.sender.balance: abort
 - (3) deduct **gasLimit** × **gasPrice** from msg.sender.balance
-
- (4) set **Gas** \leftarrow **gasLimit**
 - (5) execute Tx: deduct gas from **Gas** for each instruction
if at end (**Gas** < 0): abort, Tx is invalid (proposer keeps **gasLimit** × **gasPrice**)
 - (6) Refund **Gas** × **gasPrice** to msg.sender.balance
-
- (7) **gasUsed** \leftarrow **gasLimit** – **Gas**
 - (7a) BURN **gasUsed** × **baseFee**
 - (7b) Send **gasUsed** × (**gasPrice** – **baseFee**) to block producer



Why burn ETH ???

Recall: EIP1559 goals (informal)

- users incentivized to bid their true utility for posting Tx,
- block proposer incentivized to not create fake Tx, and
- disincentivize off chain agreements.

Suppose no burn (i.e., `baseFee` given to block producer):

⇒ in periods of low Tx volume proposer would try to increase volume by offering to refund the `baseFee` *off chain* to users.

END OF LECTURE

Next lecture: writing Solidity contracts