

CS251 Fall 2025
(cs251.stanford.edu)



Solidity

Dan Boneh

<https://docs.soliditylang.org/en/latest/>

Recap

World state: set of accounts identified by 20-byte address.

Two types of accounts:

(1) owned accounts (EOA): address = $H(pk)$

(2) contracts: address = $H(\text{CreatorAddr}, \text{CreatorNonce})$

every contract has its own storage array $S[\text{bytes32}] \rightarrow \text{bytes32}$

Recap: Transactions


- **To:** 20-byte address ($0 \rightarrow$ create new account)
- **From:** 20-byte address
- **Value:** # Wei being sent with Tx ($1 \text{ Wei} = 10^{-18} \text{ ETH}$, $1 \text{ GWei} = 10^{-9} \text{ ETH}$)
- Tx fees (EIP 1559): **gasLimit**, **maxFee**, **maxPriorityFee**
- **calldata:** what contract function to call & arguments
 - if To = 0: create new contract **code = (init, body)**
- **[signature]:** if Tx initiated by an owned account (EOA)

Recap: Blocks

Validators collect Tx from users:

- ⇒ run Tx sequentially on current world state
- ⇒ new block contains **updated world state**, Tx list, log msgs

The Ethereum blockchain: abstractly



EVM mechanics: execution environment

Write code in Solidity (or another front-end language)

- ⇒ compile to EVM bytecode, e.g., using **solc** compiler
- ⇒ validators run the contract's EVM bytecode in response to a Tx

The EVM

The EVM

see <https://www.evm.codes>

Stack machine (like Bitcoin) but with JUMP

- max stack depth = 1024
- program aborts if stack size exceeded; block proposer keeps gas

A contract can create or call another contract

- There are several ways to call another contract
- Using the CALL instruction to call another contract creates a new execution frame that is deleted on return

The EVM

see <https://www.evm.codes>

The EVM maintains three types of zero initialized memory per contract.
All three are private to the contract that owns them (e.g., nameCoin)

- **Persistent storage** (on blockchain): SLOAD, SSTORE (expensive)
- **Volatile memory** (lives for a single Tx, one per execution frame):
MLOAD, MSTORE (very cheap, 3 Gas)
- **Transient memory** (lives for a single Tx, one per contract):
TLOAD, TSTORE (cheap, 100 Gas)
- LOG0(data): write data to log (easily read by a block explorer)
- Tx calldata (16 gas/byte): read-only, readable by EVM only in current Tx

Every instruction costs gas

Why charge gas?

- Tx fees (gas) prevents submitting Tx that runs for many steps.
- During high load: block proposer chooses Tx from mempool that maximize its income.


if **gasUsed ≥ gasLimit**: block proposer keeps gas fees (from Tx originator)

calculated by EVM


specified in Tx

Gas prices spike during congestion

GasPrice in Gwei: $1.672 \text{ Gwei} = 1.672 \times 10^{-9} \text{ ETH}$



Average Tx fee in USD:



Gas calculation: EIP1559

Every block has a “baseFee”: the **minimum** gasPrice for Tx in the block

baseFee is computed from total gas in earlier blocks:

- earlier blocks at gas limit (45M gas) \Rightarrow base fee goes up 12.5%
 - earlier blocks empty \Rightarrow base fee decreases by 12.5%
- } interpolate
in between

If earlier blocks at “target size” (22.5M gas) \Rightarrow baseFee does not change

Gas calculation

A transaction specifies three parameters:

- **gasLimit**: max total gas allowed for Tx
- **maxFee**: maximum allowed gas price
- **maxPriorityFee**: additional “tip” to be paid to block proposer

Computed **gasPrice** bid (in Wei = 10^{-18} ETH):

gasPrice $\leftarrow \min(\text{maxFee}, \text{baseFee} + \text{maxPriorityFee})$

Max Tx fee: **gasLimit** \times **gasPrice**

Gas calculation (informal)

gasUsed \leftarrow gas used by Tx

Send $\text{gasUsed} \times (\text{gasPrice} - \text{baseFee})$ to block proposer

BURN $\text{gasUsed} \times \text{baseFee}$



\Rightarrow total supply of ETH can decrease

Gas calculation

- (1) if **gasPrice < baseFee**: abort
 - (2) If **gasLimit × gasPrice > msg.sender.balance**: abort
 - (3) deduct **gasLimit × gasPrice** from **msg.sender.balance**
-
- (4) set **Gas** \leftarrow **gasLimit**
 - (5) execute Tx: deduct gas from **Gas** for each instruction
if at end (**Gas < 0**): abort, Tx is invalid (proposer keeps **gasLimit × gasPrice**)
 - (6) Refund **Gas × gasPrice** to **msg.sender.balance** (leftover change)
-
- (7) **gasUsed** \leftarrow **gasLimit – Gas**
 - (7a) BURN **gasUsed × baseFee** 
 - (7b) Send **gasUsed × (gasPrice – baseFee)** to block producer

Example baseFee and effect of burn


block #	gasUsed	baseFee (Gwei)	ETH burned
23573600	10,243,205	0.12574 ↓	0.001288
23573599	19,388,322 (<22.5M)	0.12794	0.002481
23573598	26,354,665	0.12525 ↑	0.003301
23573597	23,873,377 (>22.5M)	0.12429 ↓	0.002967
23573596	21,155,077 (<22.5M)	0.12523 ↑	0.002649
23573595	44,988,950 (>22.5M)	0.11132	0.005008

≈ gasUsed × baseFee

new issuance > burn ⇒ ETH inflates

new issuance < burn ⇒ ETH deflates

Eth total supply (last 5 years)



Why burn ETH ???

EIP1559 goals (informal):

- users incentivized to bid their true utility for posting Tx,
- block proposer incentivized to not create fake Tx, and
- disincentivize off chain agreements.

Suppose no burn (i.e., baseFee given to block producer):

⇒ in periods of low Tx volume proposer would try to increase volume by offering to refund the baseFee to users.

Let's look at the Ethereum blockchain

etherscan.io:

Latest Blocks		
Bk	15778674 7 secs ago	Fee Recipient Fee Recipient: 0x6d2...766 138 txns in 12 secs
Bk	15778673 19 secs ago	Fee Recipient Lido: Execution Layer Re... 111 txns in 12 secs
Bk	15778672 31 secs ago	Fee Recipient Flashbots: Builder 313 txns in 12 secs
Bk	15778671 43 secs ago	Fee Recipient Lido: Execution Layer Re... 34 txns in 12 secs

From/to address

Tx value

From	To	Value
0x39feb77c9f90fae6196...	0x52de8d3febd3a06d3c...	0.088265 Ether
areyouguy.eth	0x404f5a67f72787a6dbd...	0.2 Ether
Optimism: State Root Pr...	Optimism: State Commit...	0 Ether
0xb3336d324ed828dbc8...	Uniswap V3: Router 2	0 Ether
0x1deaf9880c1180b023...	Uniswap V3: Router 2	0.14 Ether
0x10c5a61426b506dcba...	Uniswap V2: Router 2	0 Ether
defiantplatform.eth	0x617dee16b86534a5d7...	0 Ether

Let's look at a transaction ...

Transaction ID: 0x14b1a03534ce3c460b022185b4 ...

From: [0x1deaf9880c1180b02307e940c1e8ef936e504b6a](#)

To: Contract [0x68b3465833fb72a70ecdf485e0e4c7bd8665fc45](#)
(Uniswap V3: Router 2)

Value: 0.14 Ether (\$182)

Data: Function: [multicall\(\)](#) [calls multiple methods in a single call]

Contract generated a call to Contract 0xC02aaA39b22 ... (value:0.14)

Let's look at the To contract ...

Contract 0xC02aaA39b223FE8D0A0e5C4F27eAD9083C756Cc2

(Wrapped ETH: called from Uniswap V3: Router 2)

Balance: **4,133,236 Ether**

Code: 81 lines of solidity

} anyone can read

```
function withdraw(uint wad) public {
    require(balanceOf[msg.sender] >= wad);
    balanceOf[msg.sender] -= wad;
    msg.sender.transfer(wad);
    Withdrawal(msg.sender, wad); // emit log event
}
```

code snippet

Remember: contracts cannot keep secrets!

Contract 0xC02aaA39b223FE8D0A0e5C4F27eAD9083C756Cc2
(Wrapped ETH)

etherscan.io

Code Read Contract Write Contract
(storage) (see API)

Read Contract Information

1. name
Wrapped Ether *string*

2. totalSupply
4133296938185062975508724 *uint256*

Anyone can read contract state in storage array
⇒ never store secrets in contract!

Solidity variables stored in S[] array

Solidity

docs: <https://docs.soliditylang.org/en/latest/>

Several IDE's available

Contract structure

```
interface IERC20 {  
    function transfer(address _to, uint256 _value) external returns (bool);  
    function totalSupply() external view returns (uint256);  
    ...  
}  
  
contract ERC20 is IERC20 {      // inheritance  
    address owner;  
    constructor() public { owner = msg.sender; }  
    function transfer(address _to, uint256 _value) external returns (bool) {  
        ... implementation ...  
    }  }
```

Value types

- uint256
- address (bytes20) // address is a 20-byte value
 - `_address.balance`, `_address.send(value)`, `_address.transfer(value)`
 - call: send Tx to another contract
`(bool success,) = _address.call{value: msg.value/2, gas: 1000}(args);`
 - delegatecall: load code from another contract into current context
- bytes32
- bool

Reference types

- structs
- arrays
- bytes
- strings
- mappings:
 - Declaration: mapping (address => uint256) **balances**;
 - Assignment: balances[addr] = value;

```
struct Person {  
    uint128 age;  
    uint128 balance;  
    address addr;  
}  
Person[10] public people;
```

Globally available variables

- **block**: .blockhash, .coinbase, .gaslimit, .number, .timestamp
- **gasLeft()**
- **msg**: .data, .sender, .sig, .value
- **tx**: .gasprice, .origin
- abi: encode, encodePacked, encodeWithSelector, encodeWithSignature
- Keccak256(bytes), sha256(bytes)
- **require, assert** e.g.: require(msg.value > 100, "insufficient funds sent")

A → B → C → D:
at D: msg.sender == C
tx.origin == A

Function visibilities

- **external**: function can only be called from outside contract.
Arguments read from calldata
- **public**: function can be called externally and internally.
`function foo(bytes memory data) public {} // data is always copied from calldata to memory`
`function foo(bytes data) public {} // data is not copied to memory, therefore read-only`
- **private**: only visible inside contract
- **internal**: only visible in this contract and contracts deriving from it
- **view**: only read storage (no writes to storage)
- **pure**: does not touch storage

```
function f(uint a) private pure returns (uint b) { return a + 1; }
```

Inheritance

- Inheritance

```
contract Owned {  
    address owner;  
    constructor() { owner = msg.sender; }  
    modifier onlyOwner {  
        require(msg.sender == owner); _; } }
```

```
contract Destructable is Owned {  
    function destroy() public onlyOwner { selfdestruct(owner); } }
```

code of contract “owned” is compiled into contract Destructable

- Libraries: library code is executed in the context of calling contract
 - library **Search** { function **IndexOf()**; }
 - contract A { function B { **Search.IndexOf()**; } }

ERC20 tokens

- <https://github.com/ethereum/EIPs/blob/master/EIPS/eip-20.md>
- A standard API for fungible tokens that provides basic functionality to transfer tokens or allow the tokens to be spent by a third party.
- An ERC20 token is itself a smart contract that maintains all user balances:
`mapping(address => uint256) internal balances;`
- A standard interface allows other contracts to interact with every ERC20 token.
No need for special logic for each token.

ERC20 token interface

- function **transfer**(address _to, uint256 _value) external returns (bool);
- function **transferFrom**(address _from, address _to, uint256 _value) external returns (bool);
- function **approve**(address _spender, uint256 _value) external returns (bool);
- function **totalSupply()** external view returns (uint256);
- function **balanceOf**(address _owner) external view returns (uint256);
- function **allowance**(address _owner, address _spender) external view returns (uint256);

How are ERC20 tokens transferred?

```
contract ERC20 is IERC20 {  
  
    mapping (address => uint256) internal balances;  
  
    function transfer(address _to, uint256 _value) external returns (bool) {  
        require(balances[msg.sender] >= _value, "ERC20_INSUFFICIENT_BALANCE");  
        balances[msg.sender] -= _value;  
        balances[_to] += _value;  
        emit Transfer(msg.sender, _to, _value); // write log message  
        return true;  
    }  
}
```

Tokens can be minted by a special function **mint(address _to, uint256 _value)**

Calling other contracts

- Addresses can be cast to contract types:

```
address _token = 0x2b34aced3456781243651234de348791;
```

```
ERC20Token tokenContract = ERC20Token(_token); // type cast
```

- When calling a function on an external contract, Solidity will automatically handle ABI encoding, copying to memory, and copying return values.

```
// call the `transfer` function at address tokenContract  
(bool success,) = tokenContract.transfer(_to, _value);
```

this causes the EVM to send a message from origin contract to tokenContract.

ABI encoding and decoding

When calling a contract, the **calldata** in the transaction is an ABI encoding of:

- (1) First 4 bytes of calldata: a **function selector** indicating what function in the contract to run
 - The function selector is the first 4 bytes of the hash of the function signature:
for `transfer`, this looks like `bytes4(keccak256("transfer(address,uint256)"))`;
- (2) The rest of calldata is an ABI encoding of the function arguments

Contracts can also have two special ``last resort'' functions: **receive** and **fallback**

- These functions are called if no function in the contract matches the selector in calldata

```
receive() external payable { code }    // called if calldata is empty (i.e., pure ETH transfer)
```

```
fallback() external payable { code }   // called if calldata is not empty
```

An example ABI encoding

```
ERC20Token tokenContract = ERC20Token(_token);  
address _to;
```

The easy way to call the transfer function:

```
// Solidity compiler creates the ABI encoding  
(bool success,) = tokenContract.transfer( _to, 1 ether);
```

the function selector
for `transfer`

Or you can do the ABI encoding yourself:

```
action = abi.encodeWithSelector( tokenContract.transfer.selector, _to, 1 ether);  
(bool success,) = tokenContract.call(action);
```

Stack variables

- Stack variables generally cost the least gas
 - can be used for any simple types (anything that is <= 32 bytes).
 - `uint256 a = 123;`
 - All simple types are represented as `bytes32` at the EVM level.
 - Only 16 stack variables can exist within a single scope.

Calldata

- Calldata is a read-only byte array.
- Every byte of a transaction's calldata costs gas
(16 gas per non-zero byte, 4 gas per zero byte).
- It is cheaper to load variables directly from calldata, rather than copying them to memory.
 - This can be done by marking a function as `external` or `public` without marking the argument as `memory`

Memory (compiled to MSTORE, MLOAD)

- Memory is a byte array.
- Complex types (anything > 32 bytes such as structs, arrays, and strings) must be stored in memory or in storage.

```
string memory name = "Alice";
```

- Memory is cheap, but the cost of memory grows quadratically.

Storage array (compiled to SSTORE, SLOAD)

- Using storage is very expensive and should be used sparingly.
- Writing to storage is most expensive.
Reading from storage is cheaper, but still relatively expensive.
- mappings and state variables are always in storage.
- Some gas is refunded when storage is deleted or set to 0
- Trick for saving gas: variables < 32 bytes can be packed into 32 byte slots.

Event logs

- Event logs are a cheap way of storing data that does not need to be accessed by any contracts.
- Events are stored in transaction receipts, rather than in storage.

Security considerations

- Are we checking math calculations for overflows and underflows?
 - done by the compiler since Solidity 0.8.
- What assertions should be made about function inputs, return values, and contract state?
- Who is allowed to call each function?
- Are we making any assumptions about the functionality of external contracts that are being called?

Reentrancy bugs

```
contract Bank {  
    mapping(address => uint256) public userBalances;  
  
    function getUserBalance(address user) constant public returns(uint256) {  
        return userBalances[user];    }  
  
    function addToBalance() public payable { // customer deposits funds  
        userBalances[msg.sender] = userBalances[msg.sender] + msg.value;    }  
  
    function withdrawBalance() public { // customer withdraws its entire balance  
        uint256 amountToWithdraw = userBalances[msg.sender];  
        // send ETH from Bank contract to caller ... vulnerable!  
        (bool succ,) = msg.sender.call{value: amountToWithdraw}("");  
        require(succ, "Withdraw failed");  
        userBalances[msg.sender] = 0; // if success, clear user's balance  
    }  
}
```

```
contract Attacker {  
    Bank bank;  
  
    constructor(address bankAddress) payable {  
        bank = Bank(bankAddress);  
        bank.addToBalance{value: 75}(); // Deposit 75 Wei  
    }  
  
    function triggerAttack() public { bank.withdrawBalance(); }  
  
    receive() external payable {  
        // repeat as long as we have enough gas and the bank still has funds  
        if (gasleft() > 10000 && address(bank).balance >= 75) {  
            bank.withdrawBalance();  
        }  
    }  
}
```

ETH balance of Bank contract

Why is this an attack?

step 1: Attacker → Bank.addToBalance(75)

step 2: Attacker → Bank.withdrawBalance →
Attacker.receive → Bank.withdrawBalance →
Attacker.receive → Bank.withdrawBalance → ...

Withdraw 75 Wei from Bank contract at each recursive step !!

- Need to ensure overall transaction does not fail; ensured by if

How to fix: method 1

```
function withdrawBalance() public {
    uint amountToWithdraw = userBalances[msg.sender];

    // checks
    require(amountToWithdraw > 0, "No balance to withdraw");

    // effects
    userBalances[msg.sender] = 0;      // clear user's balance

    // interactions
    (bool success, ) = msg.sender.call{value:amountToWithdraw}("");
    require(success);      // revert transaction on failure
}
```

How to fix: method 2

```
bool transient locked; // a flag in the contract's transient memory
```

```
modifier nonReentrant {
    require(!locked, "Reentrancy attempt"); // function prologue
    locked = true;
    ;
    locked = false; // reset guard in function epilogue
}
```

```
function withdrawBalance() public nonReentrant { // a protected function
    ...
}
```

END OF LECTURE

Next lecture: DeFi contracts